① Aligned signal: [X] = France (correct)

I	went	to	France	and	visited	the	Eiffel	Tower	.
0.93	0.81	0.91	0.41	0.93	0.62	0.55	**0.97**	0.14	0.35
			top: Paris .52					*top: iffel .42*	

② Null signal: [X] = [M] (mask)

I	went	to	[M]	and	visited	the	Eiffel	Tower	.
0.81	0.70	0.61	—	0.80	0.57	0.78	**0.82**	0.39	0.44
			top: Paris .99						

Mask = neutral. Correct prediction preserved, confidence drops 0.97→0.82.

③ Adversarial (unrelated): [X] = banana

I	went	to	banana	and	visited	the	Eiffel	Tower	.
0.77	0.35	0.26	0.01	0.64	0.41	0.62	**0.33**	0.43	0.39
	top: a .34	*top: a .34*	*top: Paris .37*				*London .06*		

Nonsense disrupts all predictions. Target scatters: Eiffel 0.33, London 0.06, Empire 0.04.

④ Adversarial (related): [X] = Japan

I	went	to	Japan	and	visited	the	Tokyo ✗	Tower	.
0.87	0.71	0.83	0.74	0.92	0.71	0.35	**0.91**	0.28	0.36
						top: a .39			

Semantically related error → model confidently predicts wrong answer (Tokyo 0.91).

Context signal hierarchy (Remark 1)

Aligned (France → Eiffel **0.97**) > **Null** ([M] → Eiffel **0.82**) > **Adversarial** (Japan → Tokyo 0.91 ✗)
Mask tokens are neutral: they reduce confidence but preserve the correct prediction direction.
Erroneous tokens actively bias predictions toward wrong answers — especially when semantically related.
T2M converts adversarial → null (safe); T2T may replace adversarial → adversarial (unsafe).

Correct token | Erroneous token | Mask [M] | Predicted (correct) | Predicted (degraded) | Predicted (wrong)